

Distributed Systems Paradigms

José Orlando Pereira

Departamento de Informática
Universidade do Minho



Motivation

- Handle a large number of clients and requests with a single server
- The “c10k problem” in 1999:
 - <http://www.kegel.com/c10k.html>
- Examples:
 - financial, games, ...
 - notifications in mobile apps
 - machine-to-machine (M2M)

Case study

- Simple chat server:
 - Forward all messages to all clients
- Simplest form of multi-user server-based application:
 - Input from clients
 - Internal state and logic
 - Output to clients triggered by input
- Consider:
 - Large number of clients
 - Slow connections



First threaded solution

- For each connection:
 - Handler thread
- When reading, write to all other connections
- Use buffering:
 - At user level (streams): To minimize system calls
 - In the kernel (socket): To cope with slow readers


Sockets in java.net

```
ServerSocket ss=new ServerSocket(12345);  
  
while(true) {  
    Socket s=ss.accept();  
  
    // start handler thread  
  
    s.close();  
}
```

Buffers in java.net

```
InputStream is=new BufferedInputStream(s.getInputStream());
OutputStream os=new BufferedOutputStream(s.getOutputStream());
while(true) {
    // i/o
    is.read(...)
    for(var os: connections) {
        os.write(...);
    }
}
```

First threaded solution




Buffers in java.net

```
InputStream is=new BufferedInputStream(s.getInputStream());  
OutputStream os=new BufferedOutputStream(s.getOutputStream());  
while(true) {  
    // i/o  
    is.read(...)  
    for(var os: connections) {  
        os.write(...);  
        os.flush();  
    }  
}
```



What if flush blocks?

Writer thread



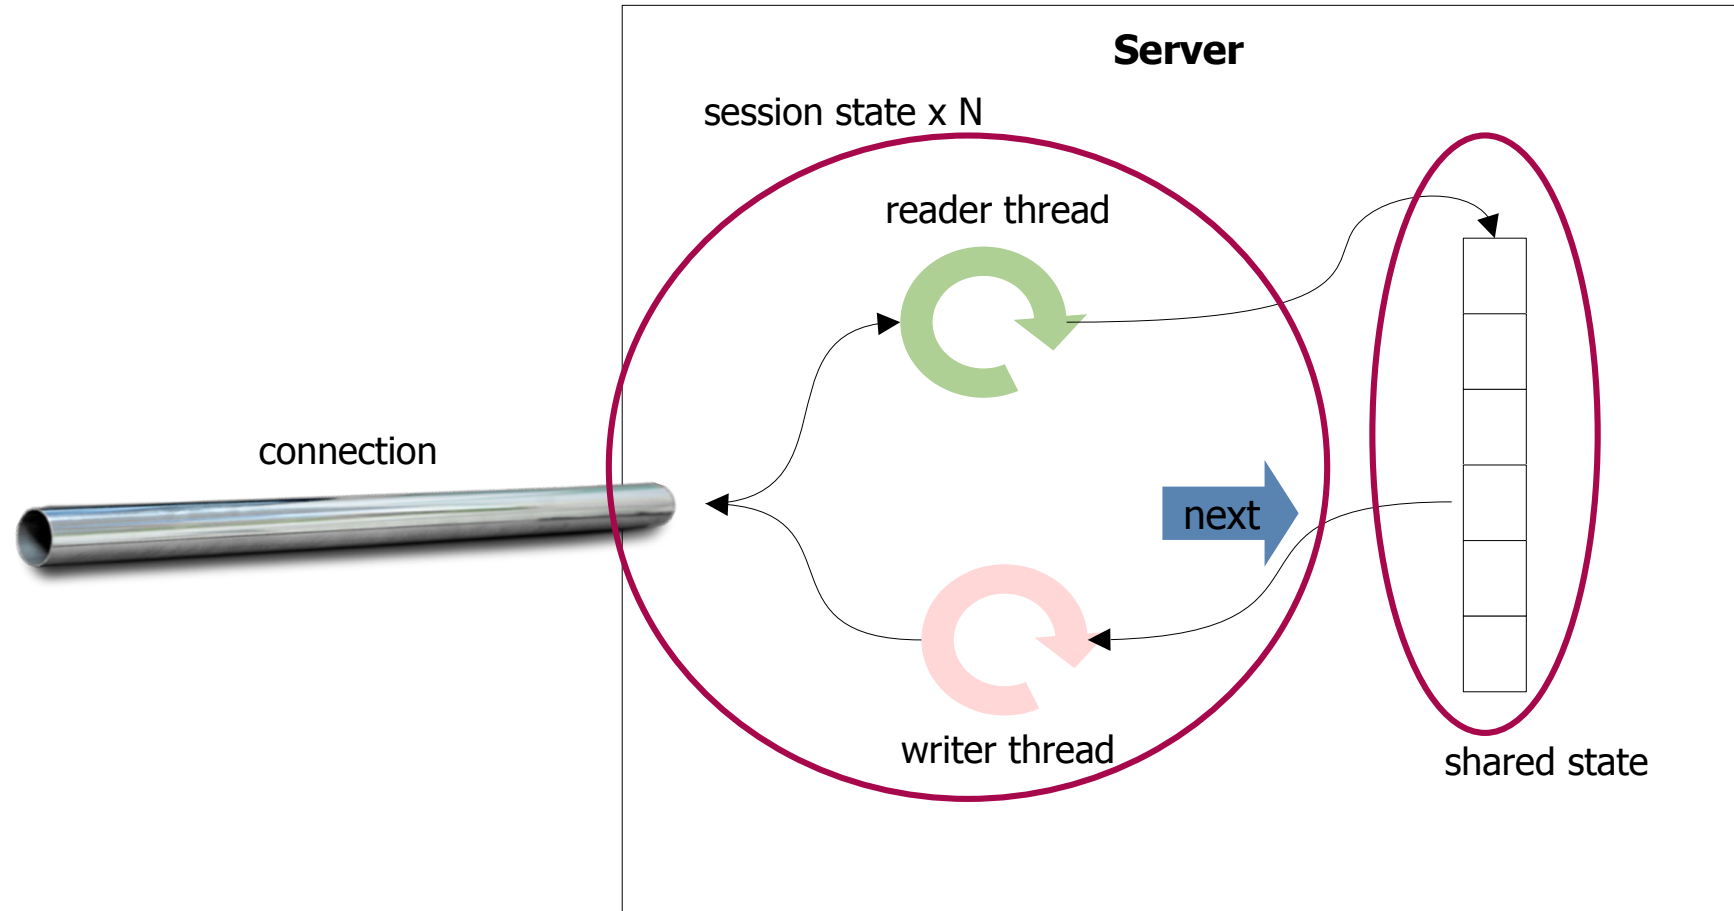
Memory

- Memory: n connections x messages in transit ($\sim n^2$)
 - Caused by data copying in stacked abstractions
 - Serialization!
 - Overhead in allocation and garbage collection

Second threaded solution

- Keep a shared queue
- For each connection:
 - Reader thread + Pointer into queue + Writer thread
- When reading, insert in shared queues and notify writer threads
- When writing, advance pointer
- Lazily, remove prefix from shared queue

Server architecture



Sockets in java.nio

```
ServerSocketChannel ss=ServerSocketChannel.open();  
ss.bind(new InetSocketAddress(12345));  
  
while(true) {  
    SocketChannel s=ss.accept();  
  
    // start i/o threads  
}
```

Sockets in java.nio

Explicit allocation of
byte buffers

```
ByteBuffer buf=ByteBuffer.allocate(100);
```

```
s.read(buf);  
buf.flip();
```

```
for(SocketChannel r: receivers) {  
    r.write(buf.duplicate());  
}  
buf.clear();
```

Shared buffers


- Memory used: messages in transit ($\sim n$)
- Ideally, never allocate or dispose of memory in normal operation:
 - No overhead, but...
 - Needs reference counting to know when to reuse
- Good tradeoff:
 - Take advantage of the garbage collector
 - Reduce the copying and duplication for large objects

Threads summary

- Simple programming model
- Problems:
 - Memory overhead (stacks)
 - Context switches and lock contention
 - “Thundering herd”, hidden queues, and fairness


Buffers in java.nio

- Buffer = Array + Indexes:



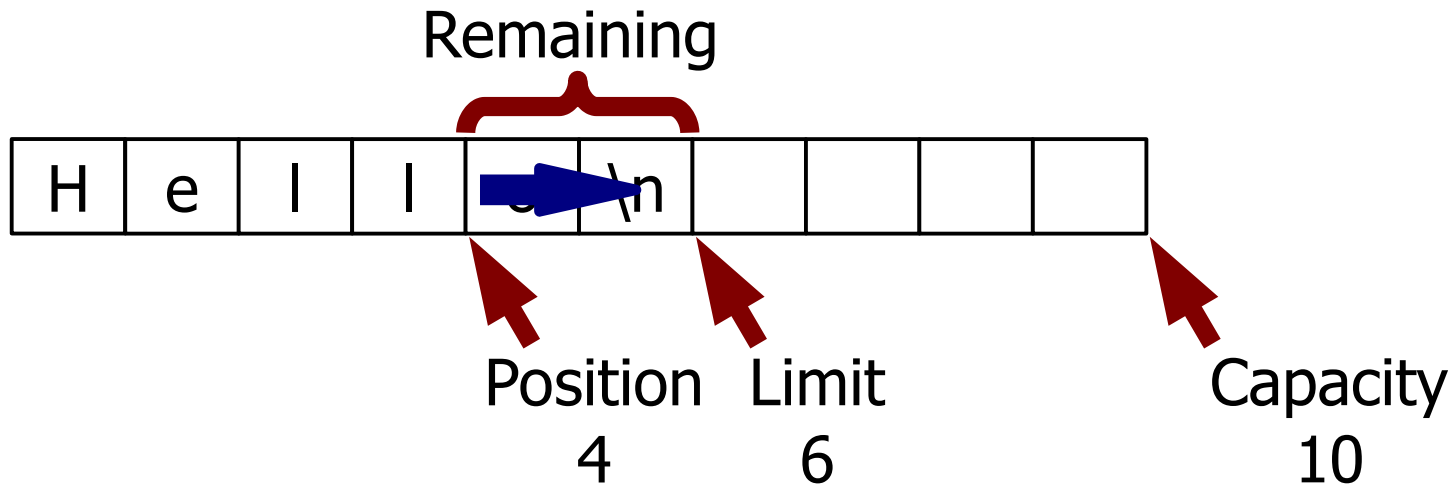
Buffers in java.nio

- Put/read: advances position, sets content




Buffers in java.nio

- Get/write: advances position, gets content

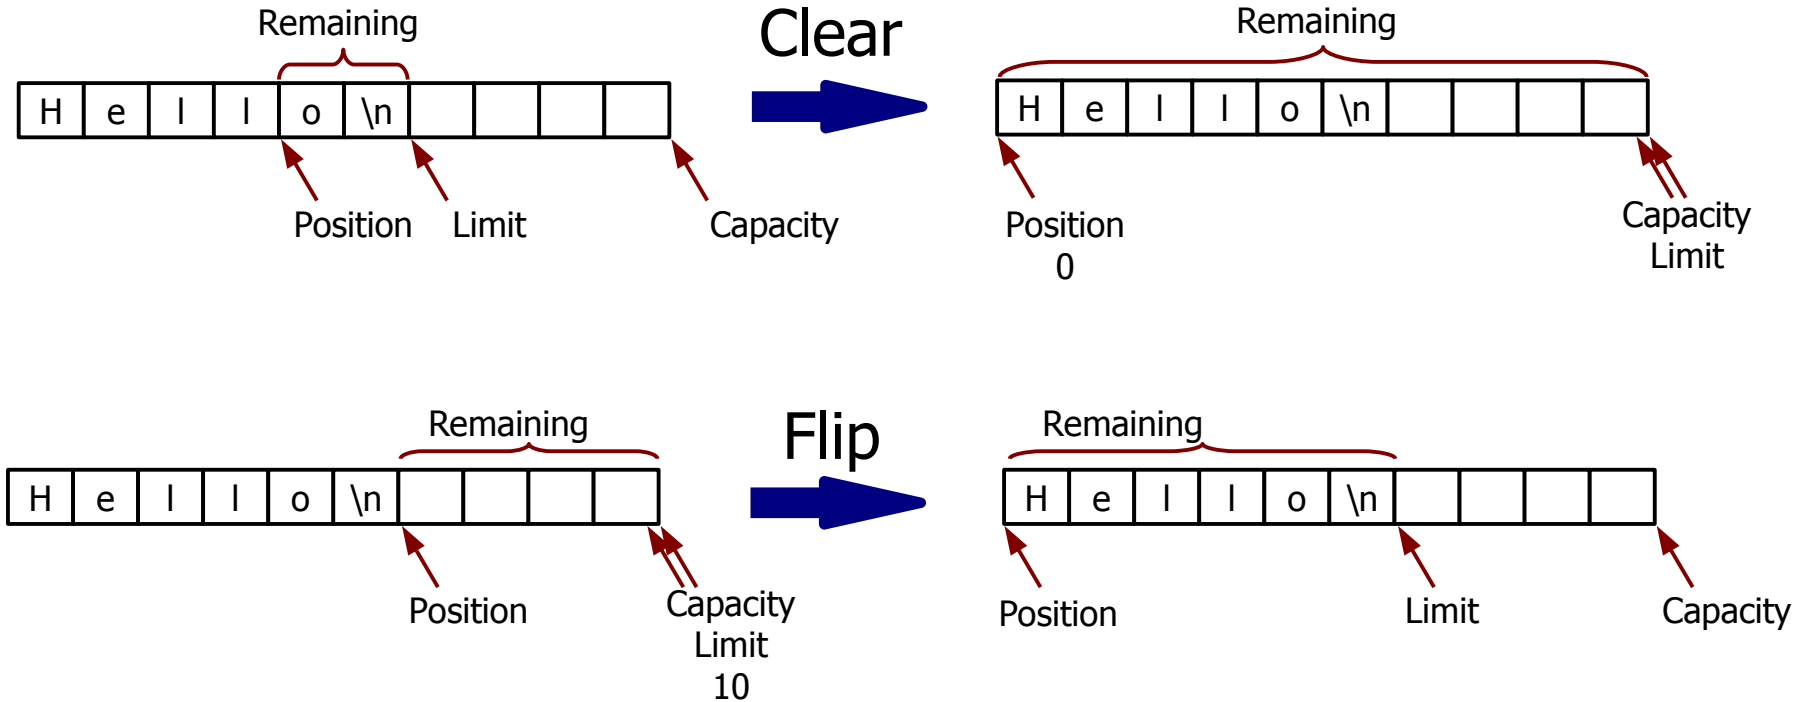


Buffers in java.nio

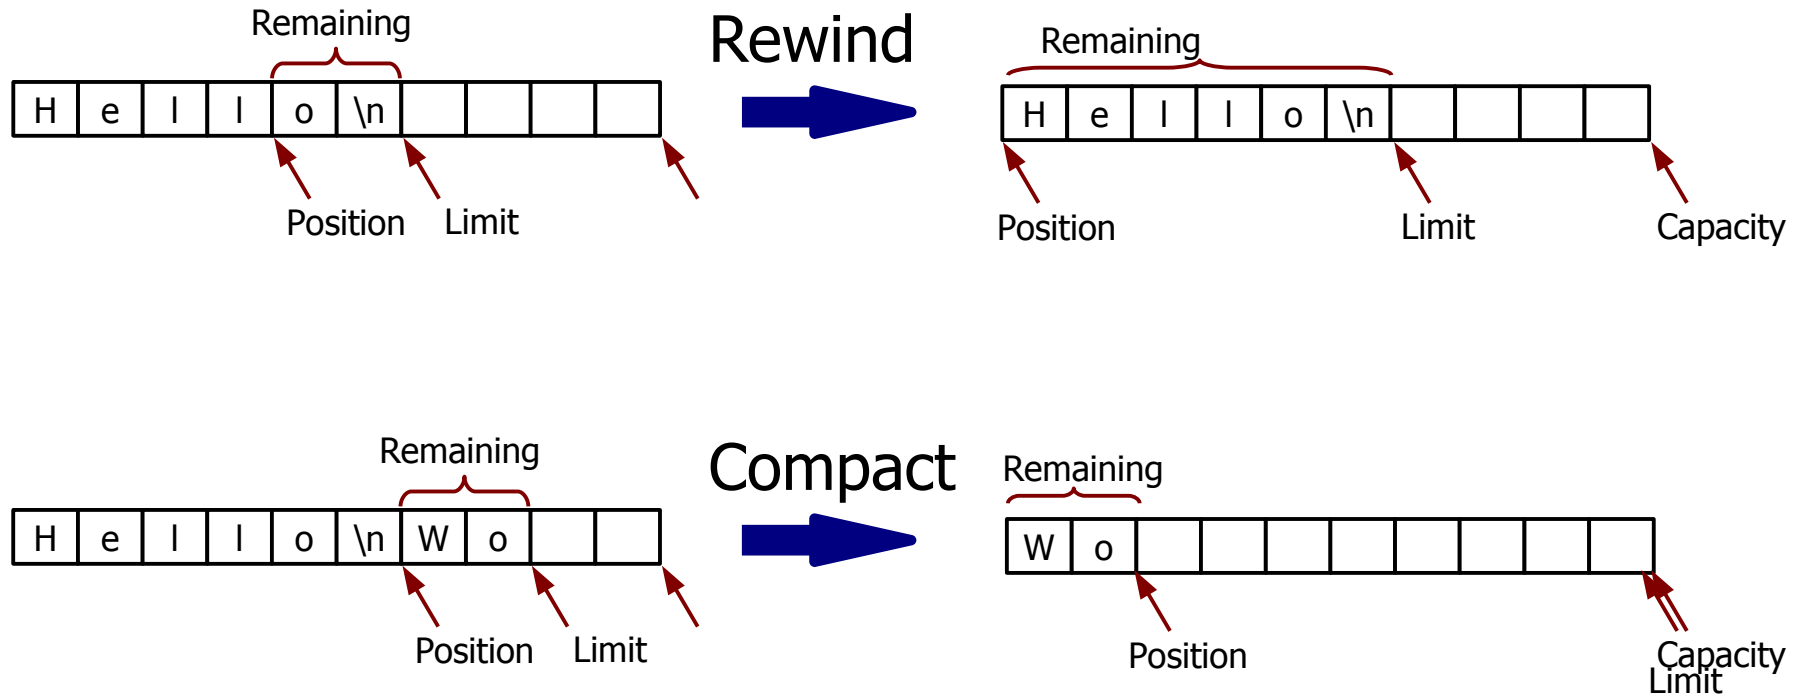
- Duplicate: multiple pointers into the same array



Buffers in java.nio



Buffers in java.nio



Serialization

- Strings need to be encoded to UTF-8, ASCII, ... and decoded back when received

```
Charset utf8 = StandardCharsets.UTF8;  
String s = ...  
ByteBuffer b = ...  
  
b = utf8.encode(s);  
  
s = utf8.decode(b);
```

Serialization

- Primitive types can be translated to and from bytes with `getInt/putInt`, `getFloat/putFloat`, etc...
- They are represented according to the current byte order of the byte buffer
- The byte order is “big endian” by default but can be changed